

JOSHUA PIÑA

Columbus, GA | 954-470-1449 | joshuadariuspina@gmail.com
[linkedin.com/in/joshuadpina](https://www.linkedin.com/in/joshuadpina) | joshuapina.github.io
Active Secret Clearance | U.S. Army Veteran

Professional Summary

Data Scientist with an Active Secret Clearance and 9+ years of U.S. Army and GovCon-related service. Hands-on production experience designing ETL pipelines, relational database infrastructure, prediction models, and applications that utilize full MLOps. Built and deployed a cross-platform data pipeline reconciling structurally incompatible multi-source exports into a cloud PostgreSQL backend. Proficient in Python and SQL; experience bridging Data Engineering and Data Science workflows. Military-trained to operate with precision and accountability in mission-critical environments. Recent B.S. in Data Science (GSU, Dec 2025); open to relocation.

Technical Skills

Languages: Python, SQL, R

Big Data & Cloud: Git, Docker, Hadoop, PySpark, AWS

ML & AI: PyTorch, scikit-learn, XGBoost, CatBoost, LightGBM, Prophet, Transformers, WandB, Optuna

Data Engineering: ETL Pipeline Design, PostgreSQL, PostGIS, pgvector, Supabase, NoSQL

APIs & Integration: REST APIs, FastAPI, GoToWebinar API, Supabase API, Gemini API

Security & Auth: Row Level Security (RLS), OAuth2

Pipeline Observability: Structured Logging (Rich), JSON Manifests, Data Lineage (CCDS, DVS, Git), PyTest

Projects

PCI Attendee Database — Production ETL & BI Platform

Dec 2025 – Mar 2026

Data Engineer: Python, PostgreSQL, Supabase, Streamlit, Rich, RLS, OAuth2

- Cross-Platform ETL: Reconciled structurally incompatible exports from GoToWebinar and Microsoft Teams against external CSV reference tables, unifying attendance, schedules, and poll responses into a PostgreSQL relational database; directly analogous to integrating heterogeneous C2 data sources.
- SQL Validation & QA: Authored post-ingestion SQL validation queries to verify record counts against source files, resolve cross-platform duplicate entries, and flag nulls in critical fields prior to every reporting cycle.
- Pipeline Observability: Implemented a JSON run manifest and structured Rich-based logging tracking ingestion counts, fuzzy-match rates, and errors, full auditability across every pipeline execution.
- BI Delivery: Delivered a Supabase-backed Streamlit dashboard and PowerBI reports embedded in Microsoft Teams; reduced monthly reporting lead time by >90%.
- Security: Architected RLS and OAuth2 from the start, ensuring data isolation and secure access for non-technical staff.

Spatiotemporal Burglary Risk Forecasting

Aug – Dec 2025

Lead Data Scientist & PM: Python, PostGIS, WandB, Docker, PyTest, GitHub Projects

- Led a 4-person team building a modular end-to-end ETL and ML pipeline integrating five external sources: REST APIs (Open-Meteo), government open data (Atlanta PD, FBI NIBRS), U.S. Census records, and GIS shapefiles into a unified spatiotemporal modeling framework.
- Created distance-based geospatial features using PostGIS; integrated weather, demographic, and temporal features into a unified observation schema (1M+ observations across 117,749 incidents).
- Containerized the full pipeline and Streamlit dashboard with Docker for reproducible, environment-agnostic deployment; published benchmark dataset on Kaggle.

Ticket Heroes

Aug – Dec 2025

Lead Data Scientist: Lambda Architecture, Python, Streamlit, Spotify API, Ticketmaster API

- Built a serverless AWS Lambda ingestion pipeline collecting real-time concert ticket pricing data, constructing two multi-year panel datasets enriched with Spotify popularity metrics, venue characteristics, and Ticketmaster event data.

- Applied time series forecasting and ML regression to predict optimal purchase timing; deployed findings via an interactive Streamlit recommendation dashboard.

The Living Library — Full-Stack Research Application

Aug – Dec 2025

Full-Stack Data Engineer: FastAPI, PostgreSQL, pgvector, OAuth2, RLS, Supabase

- Built a FastAPI backend with async SQLAlchemy ORM integrating Supabase and pgvector for PostgreSQL-backed semantic search across 55+ technical textbooks via Sentence Transformer embeddings.
- Architected multi-user authentication and data isolation using OAuth2 and Row Level Security (RLS).

ML Study: Head & Neck Squamous Cell Carcinoma

Jun – Aug 2025

Lead Data Scientist & Technical PM: PyTorch (CNN, U-Net), Python, Pandas, NumPy, SciPy

- Designed and executed an extensive ETL pipeline to construct analysis-ready data panels from 1,500+ raw, unlabeled axial CT scans; achieved 90%+ dataset compression from 6GB to under 500MB via custom windowing and runtime optimization techniques.
- Applied weighted loss functions and geometric transforms to resolve class imbalance across voxel-level segmentation labels; led a multi-modal CNN and U-Net modeling study on the processed dataset.

Experience

Program Manager / Data Engineer, Public Contracting Institute

May 2022 – Present

ETL Pipelines, Database Architecture, BI Delivery, Program Management

- Independently identified a complete absence of queryable data infrastructure; proposed, scoped, and delivered a full pipeline modernization initiative without a formal mandate or dedicated engineering team.
- Engineered Nimbus2000, a production ETL pipeline reconciling structurally incompatible multi-source exports into a unified PostgreSQL relational database; implemented structured Rich-based logging and JSON run manifests for full auditability.
- Reduced monthly org-level reporting lead time by >90%; delivered secure self-service KPI access for non-technical stakeholders via Streamlit and PowerBI.

Healthcare Specialist / NCO (SGT), United States Army

Jan 2017 – Apr 2022

Secret Clearance, Leadership, Logistics, Data Management

- Maintained Active Secret Clearance throughout service; operated under strict data classification and access control requirements.
- Distilled complex medical and operational data into actionable command-level reports across 5,000+ personnel records, directly transferable to C2 reporting and operational data delivery.
- Managed \$100K+ in medical equipment inventory through accurate auditing, reporting, and logistics coordination; as an NCO, led soldiers in high-stakes, time-constrained environments.

Education

Bachelor of Science in Data Science, Georgia State University

Dec 2025

Dean's List: Spring '24, Fall '24, Summer '25 | Relevant Coursework: DBMS, ML Data Mining, DS Capstone

Certifications: Google: AI Professional | IBM: Data Analyst | CodePath: Foundations of AI Engineering

In-Progress: AWS: Solutions Architect | CompTIA: Security+

Leadership: CodePath: Tech Mentor (Aug–Dec '25) | FourBlock: Vet Professional Development Cohort (Sept–Nov '25)